■ Business Case: Jamboree Education - Linear Regression

anoeres

Content

- 1) About Jamboree
- 2) How can you help here?
- 3) Concept Used
- 4) Libraries
- 5) Exploring the data...
- 6) Exploratory data analysis
 - 6.1) Non Graphical analysis
 - 6.2) Univariate Analysis
 - 6.3) Bivariate Analysis
- 7) Data preprocessing
 - 7.1) Check for duplicate records
 - 7.2) Check for missing values
 - 7.3) Check for outlier values
 - o 7.3.1) Visualization Method
 - o 7.3.2) Z-Score Method
 - o 7.3.3) IQR Method
- 8) Prepare the data for modeling
 - 8.1) Encoding
 - 8.2) Train-Test-Split
 - 8.3) Feature scaling
- 9) Model Building
 - 9.1) Linear Regression implementation
 - 9.2) Linear Regression using OLS
 - 9.3) Test the assumptions of linear regression
 - 9.3.1) Multicollinearity check by VIF score
 - o 9.3.2) Mean of residuals
 - o 9.3.3) Linear relationship between independent & dependent variables
 - o 9.3.4) Test for Homoscedasticity
 - 9.3.5) Normality of residuals
 - 9.4) Polynomial Regression
 - o 9.4.1) Find the best degree
 - 9.4.2) Polynomial Regression implementation
 - 9.5) Regularization
 - o 9.5.1) Laso regression model
 - o 9.5.2) Ridge regression model
 - o 9.5.3) Elastic Net regression model
 - 9.6) Model Conclusion
- 10) Inference
- 11) Recommendations

1) About Jamboree

• Jamboree has **helped thousands of students like you make it to top colleges abroad.** Be it GMAT, GRE or SAT, their unique problem-solving methods ensure maximum scores with minimum effort.

• They recently launched a feature where students/learners can come to their website and check their **probability of getting into the IVY league college.** This feature estimates the chances of graduate admission from an Indian perspective.

2) How can you help here? ••

Your analysis will help Jamboree in understanding what factors are important in graduate admissions and how these factors are interrelated among themselves. It will also help predict one's chances of admission given the rest of the variables.

3) Concept Used

- Exploratory Data Analysis
- Linear Regression

4) Libraries 🐏

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from scipy.stats import shapiro
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import Lasso, Ridge, ElasticNet
import warnings
warnings.filterwarnings('ignore')
```

5) Exploring the data... 🔎

```
In [ ]: # Loading the data
         data = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Data Sets/Jamboree_admission.csv")
        data.head()
Out[]:
                                                                      LOR CGPA Research Chance of Admit
            Serial No. GRE Score TOEFL Score University Rating
                                                                 SOP
         0
                    1
                             337
                                          118
                                                                              9.65
                                                                                                         0.92
                                                                  4.5
                                                                        4.5
                    2
                             324
                                          107
                                                                  4.0
                                                                        4.5
                                                                              8.87
                                                                                                         0.76
         2
                    3
                             316
                                          104
                                                                              8.00
                                                                                           1
                                                                                                         0.72
                                                                  3.0
                                                                        3.5
                    4
                             322
                                          110
                                                                  3.5
                                                                        2.5
                                                                              8.67
                                                                                                         0.80
         4
                    5
                                          103
                                                                                           0
                                                                                                         0.65
                             314
                                                                  2.0
                                                                        3.0
                                                                              8.21
```

```
In [ ]: data.tail()
```

Dut[]:		Serial No.	GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research	Chance of Admit
	495	496	332	108	5	4.5	4.0	9.02	1	0.87
	496	497	337	117	5	5.0	5.0	9.87	1	0.96
	497	498	330	120	5	4.5	5.0	9.56	1	0.93
	498	499	312	103	4	4.0	5.0	8.43	0	0.73
	499	500	327	113	4	4.5	4.5	9.04	0	0.84

```
In []: # Checking the number of rows and columns
    print(f"The number of rows: {data.shape[0]:,} \nThe number of columns: {data.shape[1]}")

The number of rows: 500
    The number of columns: 9

In []: # Check all column names
    data.columns
```

The data set has 500 rows and 9 columns

Column Profiling:

- Serial No.: This column represents the unique row identifier for each applicant in the dataset.
- **GRE Scores:** This column contains the GRE (Graduate Record Examination) scores of the applicants, which are measured on a scale of 0 to 340
- **TOEFL Scores:** This column includes the TOEFL (Test of English as a Foreign Language) scores of the applicants, which are measured on a scale of 0 to 120.
- University Rating: This column indicates the rating or reputation of the university that the applicants are associated with.
 - The rating is based on a scale of 0 to 5, with 5 representing the highest rating.
- **SOP:** This column represents the strength of the applicant's statement of purpose, rated on a scale of 0 to 5, with 5 indicating a strong and compelling SOP.
- LOR: This column represents the strength of the applicant's letter of recommendation, rated on a scale of 0 to 5, with 5 indicating a strong and compelling LOR.
- CGPA: This column contains the undergraduate Grade Point Average (GPA) of the applicants, which is measured on a scale of 0 to 10.
- Research: This column indicates whether the applicant has research experience (1) or not (0).
- Chance of Admit: This column represents the estimated probability or chance of admission for each applicant, ranging from 0 to 1.

```
In [ ]: data.info()
       <class 'pandas.core.frame.DataFrame'>
       RangeIndex: 500 entries, 0 to 499
       Data columns (total 9 columns):
                    Non-Null Count Dtype
       # Column
          Serial No. 500 non-null int64
GRE Score 500 non-null int64
TOEFL Score 500 non-null int64
       --- -----
       0
       2
                                            int64
       3
           University Rating 500 non-null
                    500 non-null
       4
           SOP
                                            float64
                              500 non-null
       5
           LOR
                                             float64
       6
           CGPA
                            500 non-null
                                             float64
       7
           Research
                            500 non-null
                                             int64
          Chance of Admit 500 non-null
                                             float64
       dtypes: float64(4), int64(5)
       memory usage: 35.3 KB
```

6) Exploratory data analysis 📊 📈

6.1) Non Graphical analysis

```
In [ ]: # Number of unique values in each coluumn
        print("Number of unique values in each coluumn:")
        print("-" * 40)
        for elem in (data.columns):
          print(f"{elem}: {data[elem].nunique()}")
       Number of unique values in each coluumn:
       Serial No.: 500
      GRE Score: 49
       TOEFL Score: 29
       University Rating: 5
       SOP: 9
       LOR: 9
       CGPA: 184
       Research: 2
       Chance of Admit : 61
In [ ]: # Unique values in the follwing columns
        print("Unique values in the following coluumn:")
        print("-" * 39)
        req_cols = ['University Rating', 'Research']
        for elem in req_cols:
          print(f"{elem}: {data[elem].unique()}")
       Unique values in the following coluumn:
       -----
       University Rating: [4 3 2 5 1]
       Research: [1 0]
```

• The columns University Rating has 5 unique values and Research has 2 unique values so we can convert to Categorical datatype.

• Since the column Serial No is irrelevant, we can drop them.

```
In [ ]: # Convert columns to a categorical type
        data['University Rating'] = data['University Rating'].astype('category')
        data['Research'] = data['Research'].astype('category')
        # Drop the Serial No. column
        data.drop(columns=['Serial No.'], inplace=True)
In [ ]: # Checking the column names
        data.columns
Out[ ]: Index(['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR ', 'CGPA',
                'Research', 'Chance of Admit '],
               dtype='object')
In [ ]: # Removing the unwanted space in the column names
        data.rename(columns={'Chance of Admit ':'Chance of Admit'}, inplace=True)
        data.rename(columns={'LOR ':'LOR'}, inplace=True)
In [ ]: # Sanity check
        data.columns
Out[ ]: Index(['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR', 'CGPA',
                'Research', 'Chance of Admit'],
               dtype='object')
In [ ]: # Check the data type of each column
        print("Data types of each column:")
        print("-" * 27)
        data.dtypes
       Data types of each column:
Out[]: GRE Score
                                int64
        TOEFL Score
                                int64
        University Rating category
                              float64
         SOP
                              float64
        LOR
                              float64
        CGPA
         Research
                              category
         Chance of Admit
                              float64
        dtype: object
In [ ]: # Display the range of attributes
        print("Range of attributes:")
        print("-" * 20)
        data.describe(include='all').T
       Range of attributes:
Out[ ]:
                         count unique
                                        top
                                              freq
                                                       mean
                                                                   std
                                                                         min
                                                                                  25%
                                                                                         50%
                                                                                                75%
                                                                                                        max
               GRE Score
                         500.0
                                             NaN 316.47200 11.295148 290.00
                                                                              308.0000
                                                                                       317.00
                                                                                              325.00 340.00
                                  NaN
                                       NaN
             TOEFL Score
                          500.0
                                              NaN 107.19200
                                                              6.081868
                                                                        92.00
                                                                              103.0000
                                                                                       107.00
                                                                                              112.00
                                                                                                     120.00
                                  NaN
                                      NaN
        University Rating
                         500.0
                                   5.0
                                        3.0 162.0
                                                        NaN
                                                                  NaN
                                                                         NaN
                                                                                  NaN
                                                                                         NaN
                                                                                                NaN
                                                                                                        NaN
                    SOP
                         500.0
                                                     3.37400
                                                              0.991004
                                                                         1.00
                                                                                2.5000
                                                                                          3.50
                                                                                                 4.00
                                                                                                        5.00
                                  NaN NaN
                                              NaN
                    LOR
                         500.0
                                                     3.48400
                                                              0.925450
                                                                         1.00
                                                                                3.0000
                                                                                          3.50
                                                                                                 4.00
                                                                                                        5.00
                                  NaN NaN
                                             NaN
                   CGPA
                         500.0
                                  NaN NaN
                                              NaN
                                                     8.57644
                                                              0.604813
                                                                         6.80
                                                                                8.1275
                                                                                          8.56
                                                                                                 9.04
                                                                                                        9.92
                Research 500.0
                                        1.0 280.0
                                   2.0
                                                                  NaN
                                                                         NaN
                                                                                  NaN
                                                                                         NaN
                                                        NaN
                                                                                                NaN
                                                                                                       NaN
                                  NaN NaN NaN 0.72174 0.141140
                                                                                0.6300
In [ ]: # Display the statistical summary
        print("statistical summary:")
        print("-" * 20)
        data.describe().T
       statistical summary:
       -----
```

Out[]:		count	mean	std	min	25%	50%	75%	max
	GRE Score	500.0	316.47200	11.295148	290.00	308.0000	317.00	325.00	340.00
	TOEFL Score	500.0	107.19200	6.081868	92.00	103.0000	107.00	112.00	120.00
	SOP	500.0	3.37400	0.991004	1.00	2.5000	3.50	4.00	5.00
	LOR	500.0	3.48400	0.925450	1.00	3.0000	3.50	4.00	5.00
	CGPA	500.0	8.57644	0.604813	6.80	8.1275	8.56	9.04	9.92
	Chance of Admit	500.0	0.72174	0.141140	0.34	0.6300	0.72	0.82	0.97


- The GRE scores of applicants range from 290 to 340, with an average score of approximately 316.5. Most scores are clustered between 308 and 325, as indicated by the interquartile range.
- The TOEFL scores range from 92 to 120, with an average score of about 107. The majority of the scores fall between 103 and 112.
- University ratings range from 1 to 5, with the most common rating being 3
- SOP ratings vary from 1 to 5, with an average rating of about 3.37. Most ratings are between 2.5 and 4.
- LOR ratings also range from 1 to 5, with an average rating of about 3.48. The majority of ratings are between 3 and 4.
- CGPA ranges from 6.8 to 9.92, with an average CGPA of about 8.58. Most CGPAs lie between 8.13 and 9.04.
- The research experience is binary (0 or 1). Most applicants (280 out of 500) have research experience.
- The probability of admission ranges from 0.34 to 0.97, with an average of about 0.72. Most probabilities fall between 0.63 and 0.82.


In []: green_palette = ['#006400', '#008000', '#228B22', '#32CD32', '#3CB371', '#66CDAA', '#7FFF00', '#00FF7F', '#98FB98', '#ADFF2F']

6.2) Univariate Analysis **2**

```
In [ ]: # Value couts for categorical columns
        categorical_cols = data.select_dtypes(include=['category']).columns
        for elem in categorical_cols:
          print(f"Column Name: {elem}")
          print(data[elem].value_counts())
          print("_" * 35)
          print()
       Column Name: University Rating
       University Rating
       3
            162
       2
            126
            105
       5
             73
             34
       Name: count, dtype: int64
       Column Name: Research
       Research
            280
            220
       Name: count, dtype: int64
In [ ]: # Analysis of University Rating
        green_palette = ['#006400', '#008000', '#228B22', '#32CD32', '#3CB371', '#66CDAA', '#7FFF00', '#00FF7F', '#98FB98', '#ADFF2F']
        plt.figure(figsize=(14, 5))
        plt.subplot(1, 2, 1)
        label = sns.countplot(data = data, x='University Rating', hue = 'University Rating', palette = green_palette)
        for i in label.containers:
            label.bar_label(i)
        plt.subplot(1, 2, 2)
        labels = data.groupby("University Rating")["University Rating"].count().index.categories
        values = data.groupby("University Rating")["University Rating"].count().values
        plt.pie(values, labels = labels, autopct = "%1.1f%", colors = green_palette)
        plt.ylabel("count")
        plt.suptitle("Analysis of University Rating", fontsize = 15)
        plt.tight_layout()
        plt.show()
```

Analysis of University Rating





OBSERVATION

- The university rating 3 is most common amoung applicants and contribution of university rating 3 is 32%
- The university rating 2 is second most common amoung applicants and contribution of university rating 2 is 25%
- The university rating 4 is third most common amoung applicants and contribution of university rating 4 is 21%

```
In []: # Analysis of Research
palette = ['#006400', '#32CD32']


plt.figure(figsize=(14, 5))


plt.subplot(1, 2, 1)
    label = sns.countplot(data = data, x='Research', hue = 'Research', palette = palette)
    for i in label.containers:
        label.bar_label(i)

plt.subplot(1, 2, 2)
    labels = data.groupby("Research")["Research"].count().index.categories
    values = data.groupby("Research")["Research"].count().values
    plt.pie(values, labels = labels, autopct = "%1.1f%%", colors = palette)

plt.ylabel("count")
    plt.suptitle("Analysis of Research", fontsize = 15)
    plt.tight_layout()
    plt.show()
```

Analysis of Research






- Among 500 applicants, **280 applicants** have research experience.
- **56%** proportion of applicants have research experience

```
In []: # Density plot for numerical columns
numerical_df = data.select_dtypes(include=['int64','float64'])

plt.rcParams["figure.figsize"] = [15,7]
numerical_df.plot(kind="density", subplots = True, layout = (2,3), sharex = False)
plt.suptitle("Density plot for numerical columns", fontsize = 15)
plt.show()
```

Density plot for numerical columns



OBSERVATION

• The distribution of GRE scores is approximately symmetric. This suggests that the data is evenly distributed around the mean.

- The distribution of TOEFL scores shows a slight right skew, with more scores concentrated towards the lower end and a tail extending towards higher scores.
- The distribution of SOP ratings is slightly negatively skewed. There are more higher ratings indicating that most applicants have strong SOPs.
- The distribution of LOR ratings is slightly negatively skewed. There are more higher ratings suggesting that most applicants have strong LORs.
- The distribution of CGPA scores is almost symmetric, with a skewness coefficient close to zero. This indicates an even distribution around the mean CGPA.
- The distribution of chances of admission is negatively skewed, indicating that more applicants tend to have higher chances of admission.
- The distribution of chances of admission is negatively skewed. There are more applicants with higher chances of admission indicating that overall, the distribution is skewed towards higher probabilities of admission.
- Generally, most variables show a slight negative skewness, suggesting a tendency towards higher scores or ratings.

```
# kurtosis co-efficient:
In [ ]:
        print("kurtosis co-efficient")
        print("-" * 22)
        numerical_df.kurt().round(4)
       kurtosis co-efficient
                           -0.7111
Out[]: GRE Score
         TOEFL Score
                           -0.6532
         SOP
                           -0.7057
                           -0.7457
         LOR
         CGPA
                           -0.5613
                          -0.4547
         Chance of Admit
```

dtype: float64


- The distribution of GRE scores is platykurtic, indicating that it has lighter tails compared to a normal distribution. This suggests that the distribution has fewer extreme values.
- The distribution of SOP ratings is platykurtic, indicating lighter tails compared to a normal distribution. This suggests that the distribution is relatively more concentrated around the mean SOP rating.
- The distribution of LOR ratings is platykurtic, indicating lighter tails compared to a normal distribution. This suggests that extreme ratings less common.
- The distribution of CGPA scores is platykurtic, indicating lighter tails compared to a normal distribution. This suggests that the distribution is more concentrated around the mean CGPA score
- The distribution of chances of admission is platykurtic, indicating lighter tails compared to a normal distribution. This suggests that extreme chances of admission are less common.

```
In [ ]: # Box plots for numerical columns
          palette = ['#32CD32']
          plt.figure(figsize=(13, 5))
          for i, col in enumerate(numerical_cols):
               plt.subplot(1, len(numerical_cols), i+1)
               sns.boxplot(data[col],palette = palette)
               plt.title(col)
          plt.tight_layout()
          plt.show()
                   GRE Score
                                            TOEFL Score
                                                                           SOP
                                                                                                     LOR
                                                                                                                               CGPA
                                                                                                                                                   Chance of Admit
                                                                                                                                               1.0
                                                                                                                    10.0
           340
                                     120
                                                                5.0
                                                                                          5.0
                                                                                                                                               0.9
                                                                4.5
                                                                                          4.5
                                     115
           330
                                                                                          4.0
                                                                4.0
                                                                                                                                               0.8
                                                                                                                     9.0
                                                                                                                                            Chance of Admit
90 20
                                                                3.5
                                                                                          3.5
           320
                                  TOEFL Score
        GRE Score
                                                                                                                     8.5
                                                             д
3.0
                                                                                        3.0
                                     105
           310
                                                                                                                     8.0
                                                                2.5
                                                                                          2.5
                                     100
                                                                                                                                               0.5
                                                                                          2.0
                                                                2.0
           300
                                                                1.5
                                                                                          1.5
                                      95
                                                                                                                                               0.4
                                                                                                                     7.0
           290
                                                                1.0
                                                                                          1.0
                                                                                                      0
                                                                                                                                                           0
```


• Based on the absence of outliers in GRE Score, TOEFL Score, SOP, and CGPA, and the limited presence of outliers in LOR and Chance of Admit, the dataset exhibits relatively stable distributions across most numerical attributes.

6.3) Bivariate Analysis 📊 📈









- There is a strong positive (above 0.80) correlation between,
 - CGPA Vs Chance of Admit 1st Highest correlation
 - GRE Score Vs Chance of Admit 2nd Highest correlation

- TOEFL Score Vs Chance of Admit 3rd Highest correlation
- GRE Score Vs CGPA
- TOEFL Score Vs CGPA
- TOEFL Score Vs GRE Score


```
In []: # Relationship Categorical features Vs Target
green_palette = ['#006400', '#008000', '#228B22', '#32CD32', '#3CB371', '#66CDAA', '#7FFF00', '#00FF7F', '#98FB98', '#ADFF2F']

plt.figure(figsize=(15, 6))

for i, elem in enumerate(categorical_cols):
   plt.subplot(1, 2, i+1)
   sns.boxplot(x=elem, y='Chance of Admit', data=data, palette=green_palette)
   plt.title(f"{elem} vs. Chance of Admit")

plt.suptitle("Relationship Categorical features Vs Target.", fontsize = 18)
   plt.tight_layout()
   plt.show()
```


Relationship Categorical features Vs Target.



OBSERVATION

- As the University Ratings increases, the Chance of Admit is also increasing.
- If the person have research experiance, the Chance of Admit is high.

```
In []: # University Rating vs. Chance of Admit with Research
    plt.figure(figsize=(15, 5))
    sns.boxplot(x='University Rating', y='Chance of Admit', hue='Research', data=data, palette=['#228B22', '#32CD32'])
    plt.title("University Rating vs. Chance of Admit with Research")
    plt.show()
```



OBSERVATION

The difference between the chance of Admit is high for research and Non research experisnce for the University Ratung 5.

7) Data preprocessing

7.1) Check for duplicate records

```
In [ ]: # Check if there are any duplicate records
data.loc[data.duplicated()]
```

Out[]: GRE Score TOEFL Score University Rating SOP LOR CGPA Research Chance of Admit

OBSERVATION

• There is no duplicate records.

7.2) Check for missing values 🔎

```
In [ ]: # Check if there are any missing values
         data.isna().sum()
Out[]: GRE Score
                               0
         TOEFL Score
                               0
         University Rating
         SOP
         LOR
                               0
         CGPA
                               0
         Research
                               0
         Chance of Admit
         dtype: int64
In [ ]: # Null value heatmap:
         plt.figure(figsize = (16,4))
         sns.heatmap(data.isnull().T, cmap='Greens')
         plt.title('Null Values Heatmap')
         plt.show()
                                                                Null Values Heatmap
                                                                                                                                        0.100
            GRE Score
                                                                                                                                        0.075
           TOEFL Score
                                                                                                                                        0.050
       University Rating -
                                                                                                                                        0.025
                 SOP
```

OBSERVATION

LOR -

CGPA

Research

Chance of Admit

• There is no missing values.

7.3) Check for outlier values 🏂

7.3.1) Visualization Method

```
In []: # Box plots for numerical columns
palette = ['#32CD32']
plt.figure(figsize=(13, 5))
for i, col in enumerate(numerical_cols):
    plt.subplot(1, len(numerical_cols), i+1)
        sns.boxplot(data[col],color="orange")
    plt.title(col)
plt.tight_layout()
plt.show()
```


0.000

-0.025

- -0.050

- -0.075

- -0.100



7.3.2) Z-Score Method

```
In []: # Since the distribution of numerical features are almost normal for all columns we can go with Z Score
    from scipy import stats
z_score = stats.zscore(numerical_df)
    data[((z_score < -3) | (z_score > 3)).any(axis=1)]
```

Out[]: GRE Score TOEFL Score University Rating SOP LOR CGPA Research Chance of Admit

OBSERVATION

- From the above output we can say that there is no outliers using Z-Score method.
- Lets check with IQR Method as well.

7.3.3) IQR Method

```
In [ ]: # Sum of outliers in each column
        # Function to detect outliers using the IQR method
        def detect_outliers_iqr(data):
            Q1 = data.quantile(0.25)
            Q3 = data.quantile(0.75)
            IQR = Q3 - Q1
            outliers = ((data < (Q1 - 1.5 * IQR)) | (data > (Q3 + 1.5 * IQR)))
            return outliers
        # outliers for each numerical column
        outliers_gre = detect_outliers_iqr(data['GRE Score'])
        outliers_toefl = detect_outliers_iqr(data['TOEFL Score'])
        outliers_sop = detect_outliers_iqr(data['SOP'])
        outliers_lor = detect_outliers_iqr(data['LOR'])
        outliers_cgpa = detect_outliers_iqr(data['CGPA'])
        outliers_admit = detect_outliers_iqr(data['Chance of Admit'])
        # Combine outlier information into a DataFrame
        outliers_combined = pd.DataFrame({
             'GRE Score': outliers_gre,
            'TOEFL Score': outliers_toefl,
            'SOP': outliers_sop,
            'LOR': outliers_lor,
            'CGPA': outliers_cgpa,
            'Chance of Admit': outliers_admit
        })
        # Check the number of outliers in each column
        outliers_combined.sum()
Out[]: GRE Score
        TOEFL Score
                            0
        SOP
                           0
        LOR
                           1
         CGPA
                           0
         Chance of Admit
        dtype: int64
In [ ]: # Outlier value in LOR column
        data.loc[outliers_lor]['LOR']
Out[]: 347
               1.0
         Name: LOR, dtype: float64
In [ ]: # Outlier value in Chance of Admit column
        data.loc[outliers_admit]['Chance of Admit']
```

```
Out[]: 92
               0.34
        376
               0.34
        Name: Chance of Admit, dtype: float64
In [ ]: # Calculating the five point summery and Percentage of outliers for each numerical column
        def detect_outliers(column):
          # Calculating the IQR:
          Q1 = np.percentile(data[column],25)
          median = np.percentile(data[column],50)
          Q3 = np.percentile(data[column],75)
          IQR = Q3 - Q1
          # Calculating the percentage of outliers:
          lower_bound = max(data[column].min(), Q1 - 1.5 * IQR)
          upper_bound = min(data[column].max(), Q3 + 1.5 * IQR)
          outliers = data[(data[column] < lower_bound) | (data[column] > upper_bound)]
          percentage_outliers = (len(outliers)/len(data)) * 100
          print(f"For columns: {col}")
          print(f"Lower bound value: {lower_bound.round(2)} \nQ1: {Q1.round(2)} \nMedian: {median.round(2)} \nQ3: {Q3.round(2)} \nUpper
          return percentage_outliers
        # UIterating the numerical columns:
        for col in numerical_cols:
            percentage = detect_outliers(col)
            print(f"Percentage of outliers in column '{col}': {percentage:.2f}%")
            print("-" * 60)
      For columns: GRE Score
      Lower bound value: 290
      01: 308.0
      Median: 317.0
      Q3: 325.0
      Upper bound value: 340
      IQR: 17.0
      Percentage of outliers in column 'GRE Score': 0.00%
       ______
      For columns: TOEFL Score
      Lower bound value: 92
      Q1: 103.0
      Median: 107.0
      Q3: 112.0
      Upper bound value: 120
      IQR: 9.0
      Percentage of outliers in column 'TOEFL Score': 0.00%
      For columns: SOP
      Lower bound value: 1.0
      Q1: 2.5
      Median: 3.5
      Q3: 4.0
      Upper bound value: 5.0
      IQR: 1.5
      Percentage of outliers in column 'SOP': 0.00%
      For columns: LOR
      Lower bound value: 1.5
      Q1: 3.0
      Median: 3.5
      Q3: 4.0
      Upper bound value: 5.0
      IQR: 1.0
      Percentage of outliers in column 'LOR': 0.20%
      For columns: CGPA
      Lower bound value: 6.8
      Q1: 8.13
      Median: 8.56
       Q3: 9.04
      Upper bound value: 9.92
      IQR: 0.91
      Percentage of outliers in column 'CGPA': 0.00%
       ______
      For columns: Chance of Admit
      Lower bound value: 0.35
      Q1: 0.63
      Median: 0.72
      Q3: 0.82
      Upper bound value: 0.97
      IQR: 0.19
      Percentage of outliers in column 'Chance of Admit': 0.40%
```

- The column LOR has 0.20% of data as outliers but the value is in the scale of 1-5
- The column Chance of Admit has 0.40% data as outliers.
- Since there are very minimum outliers in the dataset, we are not going to treat them.

8) Prepare the data for modeling

8.1) Encoding

```
In [ ]: # Display the categorical columns
    data[categorical_cols].head()
```

Out[]:	University Rating	Research		
	0	4	1		
	1	4	1		
	2	3	1		
	3	3	1		
	1	2	0		

```
In []: # Details about the catgorical columns
    for elem in categorical_cols:
        print(f"{elem}: {data[elem].unique()}")
        print()
University Rating: [4, 3, 2, 5, 1]
```

```
University Rating: [4, 3, 2, 5, 1]
Categories (5, int64): [1, 2, 3, 4, 5]

Research: [1, 0]
Categories (2, int64): [0, 1]
```

OBSERVATION

In []: # Convert data types again to integer

• Since the data is already numerical there is no need of encoding.

```
Out[]: GRE Score
        TOEFL Score
                              int64
        University Rating
                              int64
        SOP
                             float64
        LOR
                             float64
        CGPA
                             float64
        Research
                               int64
        Chance of Admit
                             float64
        dtype: object
```

8.2) Train-Test-Split 📯

```
In []: # Lets split the data into Independent feature and dependent feature
    y = data['Chance of Admit']
    X = data.drop('Chance of Admit', axis=1)
    X.shape, y.shape

Out[]: ((500, 7), (500,))

In []: # Lets split the data into train and test
    from sklearn.model_selection import train_test_split
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In []: # Cross check
    X_train.head()
```

```
Out[ ]:
               GRE Score TOEFL Score University Rating SOP LOR CGPA Research
         249
                     321
                                  111
                                                          3.5
                                                                4.0
                                                                      8.83
         433
                     316
                                  111
                                                          4.0
                                                                5.0
                                                                      8.54
                                                                                   0
          19
                     303
                                  102
                                                          3.5
                                                                3.0
                                                                      8.50
                                                                                   0
         322
                     314
                                  107
                                                                      8.27
                                                          2.5
                                                                4.0
                                                                                   0
         332
                     308
                                  106
                                                          3.5
                                                                2.5
                                                                      8.21
```

```
In []: # Lets check the shape of train and test data
    print(f"Shape of X_train: {X_train.shape}")
    print(f"Shape of y_train: {y_train.shape}")
    print(f"Shape of X_test: {X_test.shape}")
    print(f"Shape of y_test: {y_test.shape}")

Shape of X_train: (400, 7)
    Shape of y_train: (400,)
    Shape of X_test: (100, 7)
    Shape of y_test: (100, 7)
```

8.3) Feature scaling

Lets do Standardization (Z-score Scaling) since our data exhibits a near-normal distribution.

```
In [ ]: from sklearn.preprocessing import StandardScaler
        scaler = StandardScaler()
        X_train = scaler.fit_transform(X_train) # Train data is fit and transformed
        X_test = scaler.transform(X_test) # Test data is only transformed
In [ ]: # X_train is converted to np array
        X_train
Out[]: array([[ 0.38998634, 0.6024183 , -0.09829757, ..., 0.56498381,
                 0.4150183 , 0.89543386],
                [-0.06640493, 0.6024183, 0.7754586, ..., 1.65149114,
                -0.06785154, -1.11677706],
               [-1.25302222, -0.87691722, -0.09829757, ..., -0.52152352,
                -0.13445427, -1.11677706],
               \lceil -1.34430047, -1.37002906, -1.8458099, \ldots, -1.60803084, \rceil
                -2.2157898 , -1.11677706],
               [-0.7053527, -0.38380538, -0.97205374, ..., 0.56498381,
                -1.49981038, -1.11677706],
               [-0.24896144, -0.21943477, -0.97205374, ..., 0.02173015,
                -0.55072138, -1.11677706]])
In [ ]: # Lets convert the X_train to a dataframe
        df_train = pd.DataFrame(X_train, columns = data.columns[:-1])
        df_train.head()
Out[ ]:
           GRE Score TOEFL Score University Rating
                                                       SOP
                                                                LOR
                                                                         CGPA
                                                                                Research
                                         -0.098298
            0.389986
                         0.602418
                                                   0.126796
                                                            0.564984
                                                                      0.415018
                                                                                0.895434
            -0.066405
                                         0.775459
                         0.602418
                                                   0.633979
                                                            1.651491 -0.067852 -1.116777
           -1.253022
                        -0.876917
                                         -0.098298
                                                   0.126796 -0.521524 -0.134454 -1.116777
            -0.248961
                        -0.055064
                                         -0.972054
                                                  -0.887570
                                                            0.564984 -0.517420 -1.116777
                        -0.219435
           -0.796631
                                        In [ ]: # Check the shape
        df_train.shape
Out[]: (400, 7)
In [ ]: # Lets convert the X_test to a dataframe
        df_test = pd.DataFrame(X_test, columns = data.columns[:-1])
        df_test.head()
```

Out[]:		GRE Score	TOEFL Score	University Rating	SOP	LOR	CGPA	Research
	0	1.576604	1.424271	0.775459	0.633979	0.021730	1.597217	0.895434
	1	-0.248961	0.109306	0.775459	1.141162	0.564984	0.764683	0.895434
	2	-0.157683	-0.383805	-0.972054	-1.394754	-1.064777	-1.549762	-1.116777
	3	-0.431518	0.273677	-0.098298	-0.380387	-0.521524	0.181909	-1.116777
	4	0.846378	0.766789	-0.098298	0.126796	-0.521524	0.781333	0.895434
In []:		Check the s _test.shape	•					
Out[]:	(1	00, 7)						

9) Model Building

```
9.1) Linear Regression implementation
In [ ]: # Import the required Library
        from sklearn.linear_model import LinearRegression
        # Initialize the Linear Regression model
        LR_model = LinearRegression()
        # Train the model
        LR_model.fit(X_train, y_train)
Out[]: ▼ LinearRegression
        LinearRegression()
In [ ]: # Predicting values for the test data
        y_pred_train = LR_model.predict(X_train)
        y_pred_test = LR_model.predict(X_test)
In [ ]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
        # Evaluate the model
        def evaluate_model_performance(y_true, y_forecast, model):
            # Calculate MSE
            mse = mean_squared_error(y_true, y_forecast)
            # Calculate MAE
            mae = mean_absolute_error(y_true, y_forecast)
            # Calculate RMSE
            rmse = np.sqrt(mse)
            # Calculate R-squared
            r2 = r2_score(y_true, y_forecast)
            # Number of observations
            n = len(y_true)
            # Number of predictors (features)
            p = model.n_features_in_
            # Calculate Adjusted R-squared
            adjusted_r2 = 1 - (1 - r2) * (n - 1) / (n - p - 1)
            # Return all metrics in a dictionary
            return print(f"MSE: {mse.round(4)}\nMAE: {mae.round(4)}\nRMSE: {rmse.round(4)}\nR-squared: {r2.round(2)}\nAdjusted R-squared
In [ ]: # Performance of Linear Regression
        print("Performance of Linear Regression")
        print("-"*36)
        # Metrix for train and test data
        print("Performance of Train data")
        print("-"*26)
        evaluate_model_performance(y_train, y_pred_train, LR_model)
        print("Performance of Test data")
        print("-"*26)
        evaluate_model_performance(y_test, y_pred_test, LR_model)
```


- The model shows strong performance on both the training and test data, with R-squared values of 0.82 and 0.81 respectively.
- The low MAE, MSE, and RMSE values indicate that the predictions are close to the actual values.

```
In [ ]: # Coefficients and intercept
    coefficients_df = pd.DataFrame(LR_model.coef_.reshape(1,-1), columns=data.columns[:-1])
    coefficients_df
```

```
        Out[]:
        GRE Score
        TOEFL Score
        University Rating
        SOP
        LOR
        CGPA
        Research

        0
        0.026671
        0.018226
        0.00294
        0.001788
        0.015866
        0.067581
        0.01194
```

```
In []: # Intercept
LR_model.intercept_
Out[]: 0.724174999999999
In []: # Weightage Vs Features
green_palette = ['#006400', '#008000', '#228B22', '#32CD32', '#3CB371', '#66CDAA', '#7FFF00', '#00FF7F', '#98FB98', '#ADFF2F']
plt.figure(figsize=(14, 5))
sns.barplot(x=coefficients_df.columns, y=coefficients_df.iloc[0], palette=green_palette)
plt.title('Weightage Vs Features')
plt.xlabel('Features')
plt.ylabel('Weightage')
plt.xticks(rotation=90)
plt.show()
```



OBSERVATION

• The feature CGPA has the highest weightage among all other features followed by GRE Score and TOEFL Score.

```
In [ ]: # Actual vs Predicted
plt.figure(figsize=(14, 5))


# Actual vs Predicted Plot
plt.subplot(1, 2, 1)
```

```
plt.scatter(y_train, y_pred_train)
plt.xlabel('Actual_train (y_train)')
plt.ylabel('Predicted_train (y_pred_train)')
plt.title('Training data')

# Actual vs Predicted Plot
plt.subplot(1, 2, 2)
plt.scatter(y_test, y_pred_test, color="green")
plt.xlabel('Actual_test (y_test)')
plt.ylabel('Predicted_test (y_pred_test)')
plt.title('Testing data')

plt.suptitle("Actual vs Predicted - Linear Regression", fontsize = 18)
plt.tight_layout()
plt.show()
```

Actual vs Predicted - Linear Regression



OBSERVATION

• The points are closely aligned along the diagonal line y = x, indicating that the model's predictions are fairly accurate for both training and testing data.

9.2) Linear Regression using OLS

```
In []: import statsmodels.api as sm

X_sm = sm.add_constant(X_train) # Statmodels default is without intercept, to add intercept we need to add constant.

model = sm.OLS(y_train, X_sm)
results = model.fit()

# Print the summary statistics of the model
print(results.summary())

OLS Regression Results

Dep. Variable: Chance of Admit R-squared: 0.821
Model: OLS Adj. R-squared: 0.818
```

Dep. Variable:	Chance of Admit	R-squared:	0.821
Model:	OLS	Adj. R-squared:	0.818
Method:	Least Squares	F-statistic:	257.0
Date:	Wed, 17 Jul 2024	Prob (F-statistic):	3.41e-142
Time:	06:06:53	Log-Likelihood:	561.91
No. Observations:	400	AIC:	-1108.
Df Residuals:	392	BIC:	-1076.
Df Model:	7		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]				
const x1	0.7242 0.0267	0.003 0.006	241.441 4.196	0.000 0.000	0.718 0.014	0.730 0.039				
x2 x3	0.0182 0.0029	0.006 0.005	3.174 0.611	0.002 0.541	0.007 -0.007	0.030 0.012				
x4 x5	0.0018 0.0159	0.005 0.004	0.357 3.761	0.721 0.000	-0.008 0.008	0.012 0.024				
x6 x7	0.0676 0.0119	0.006 0.004	10.444 3.231	0.000 0.001	0.055 0.005	0.080 0.019				
Omnibus:	======	======== 86.2	:======: 22	======== oin-Watson:	=======	2.050				
Prob(Omnibus):		0.0	000 Jaro	que-Bera (JB):		190.099				
Skew: Kurtosis:		-1.1 5.5		o(JB): d. No.		5.25e-42 5.65				
=======================================	======					=======				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

9.3) Test the assumptions of linear regression 🔠

9.3.1) Multicollinearity check by VIF score

VIF

Variance Inflation Factor (VIF) is a measure of the amount of multicollinearity in a set of multiple regression variables. It quantifies how much the variance of a regression coefficient is inflated due to multicollinearity with other predictors. A VIF value greater than 10 is typically considered indicative of high multicollinearity.

Interpreting VIF Values

- VIF = 1: No correlation between the predictor and other variables.
- 1 < VIF < 5: Moderate correlation.
- VIF > 5: High correlation, indicating multicollinearity, but sometimes acceptable depending on the context.
- VIF > 10: Very high correlation, indicating a serious multicollinearity problem that needs to be addressed.

```
In [ ]: # Display the train scalled data in the form of data frame
    df_train.head()
```

```
Out[ ]:
                                                                      LOR
            GRE Score TOEFL Score University Rating
                                                            SOP
                                                                                CGPA
                                                                                       Research
             0.389986
                           0.602418
                                                                  0.564984
                                                                             0.415018
                                                                                        0.895434
                                             -0.098298
                                                       0.126796
                                                       0.633979
             -0.066405
                           0.602418
                                             0.775459
                                                                  1.651491
                                                                            -0.067852 -1.116777
                          -0.876917
                                                                 -0.521524 -0.134454 -1.116777
            -1.253022
                                            -0.098298
                                                       0.126796
                                             -0.972054
                                                                  0.564984 -0.517420 -1.116777
             -0.248961
                          -0.055064
                                                       -0.887570
                                                       0.126796 -1.064777 -0.617324
             -0.796631
                          -0.219435
                                             -0.098298
```

```
In []: # Deep copy of df_train
X_t = df_train.copy()
X_t.head()
```

```
Out[]:
            GRE Score TOEFL Score University Rating
                                                           SOP
                                                                      LOR
                                                                               CGPA
                                                                                      Research
             0.389986
                           0.602418
                                            -0.098298
                                                       0.126796
                                                                 0.564984
                                                                            0.415018
                                                                                      0.895434
            -0.066405
                           0.602418
                                             0.775459
                                                       0.633979
                                                                 1.651491 -0.067852 -1.116777
            -1.253022
                          -0.876917
                                            -0.098298
                                                       0.126796 -0.521524 -0.134454 -1.116777
             -0.248961
                          -0.055064
                                            -0.972054 -0.887570
                                                                 0.564984 -0.517420 -1.116777
             -0.796631
                          -0.219435
                                            -0.098298
                                                       0.126796 -1.064777 -0.617324
```

```
In [ ]: # Import the required libraries
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Create a empty data frame
vif = pd.DataFrame()

# Create Features column
vif['Features'] = X_t.columns

# Create VIF column
vif['VIF'] = [variance_inflation_factor(X_t.values, i) for i in range(X_t.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)

# Display the VIF data frame
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

- All VIF values are below the critical threshold of 5, indicating that there is no severe multicollinearity among the predictor variables.
- The highest VIF value is 4.65 for CGPA, which is still below the acceptable range

9.3.2) Mean of residuals

The mean of the residuals in a linear regression model is a measure of how the residuals (the differences between the observed values and the predicted values) average out. In an ideal linear regression model, the mean of the residuals should be close to zero. This indicates that, on average, the predictions are accurate and the model is well-fitted.

```
In []: # Calculate residuals
    residuals_train = y_train - y_pred_train
    residuals_test = y_test - y_pred_test

# Mean of residuals
    mean_residuals_train = np.mean(residuals_train)
    mean_residuals_test = np.mean(residuals_test)

# Print the mean of residuals
    print("Mean of residuals (train):", mean_residuals_train)
    print("Mean of residuals (test):", mean_residuals_test)
```

Mean of residuals (train): 1.4419021532319221e-16 Mean of residuals (test): -0.005453623717661251

OBSERVATION

- The mean of residuals for the training data is extremely close to zero (1.44e-16). This suggests that the model has no systematic bias for the training data.
- The mean of residuals for the test data is -0.0054, which is very close to zero. This indicates that the model generalizes well to unseen data and there is no significant bias in the predictions for the test set.
- The near-zero mean of residuals for both training and test sets indicates that the linear regression model is well-fitted and unbiased.


9.3.3) Linear relationship between independent & dependent variables

Linear regression assumes that there is a linear relationship between the independent variables (predictors) and the dependent variable (response). This means that changes in the dependent variable are proportional to changes in the independent variables.

```
In []: # Linearity of variables
numerical_cols_except_target = ['GRE Score', 'TOEFL Score', 'University Rating', 'SOP', 'LOR', 'CGPA']

plt.figure(figsize=(15, 11))
for i, elem in enumerate(numerical_cols_except_target):
    plt.subplot(3, 3, i+1)
    sns.regplot(x=elem, y='Chance of Admit', data=data, color="green")
    plt.title(f"{elem} vs. Chance of Admit")

plt.suptitle("Linearity of variables", fontsize = 18)
plt.tight_layout()
plt.show()
```



- All independent variables show a positive linear relationship with the chance of admission. This means that as each variable increases, the chance of admission also tends to increase.
- The presence of these linear trends supports the assumption of linearity in the linear regression model. However, the strength of linearity varies among different predictors.

9.3.4) Test for Homoscedasticity

Homoscedasticity refers to the assumption that the variance of the errors (residuals) is constant across all levels of the independent variables. In the context of linear regression, this means that the spread of the residuals should be roughly the same for all predicted values. Homoscedasticity is important because it ensures that the model's predictions are equally reliable across all values of the independent variables. When the assumption of homoscedasticity is violated (a condition known as heteroscedasticity), the standard errors of the coefficients can be biased, leading to incorrect conclusions about the relationships between variables.


To test for homoscedasticity, there are several graphical and statistical methods that you can use:

- Residual plot: Plot the residuals against the predicted values or the independent variables. Look for any systematic patterns or trends in the spread of the residuals. If the spread appears to be consistent across all levels of the predictors, then homoscedasticity is likely met.
- Goldfeld-Quandt Test: This test is used when you suspect heteroscedasticity due to different variances in different parts of the data. It involves splitting the data into two subsets based on a specific criterion and then comparing the variances of the residuals in each subset. If the difference in variances is not significant, it suggests homoscedasticity.

It's important to note that the visual inspection of plots is often the first step to identify potential violations of homoscedasticity. Statistical tests can provide additional evidence, but they may have assumptions or limitations that need to be considered.

```
In [ ]: # Assuming y_train, y_pred_train, y_test, and y_pred_test are already defined
        plt.figure(figsize=(14, 5))
        # Residual Plot for train data
        plt.subplot(1, 2, 1)
        loss_train = y_train - y_pred_train
        sns.residplot(x=y_pred_train, y=loss_train, lowess=True, line_kws={'color': 'red'})
        plt.xlabel('Predicted_train (y_pred_train)')
        plt.ylabel('Residuals_train')
        plt.title('Residual Plot for train data', fontsize=15)
        # Residual Plot for test data
        plt.subplot(1, 2, 2)
        loss_test = y_test - y_pred_test
        sns.residplot(x=y_pred_test, y=loss_test, lowess=True, color="green", line_kws={'color': 'red'})
        plt.xlabel('Predicted_test (y_pred_test)')
        plt.ylabel('Residuals_test')
        plt.title('Residual Plot for test data', fontsize=15)
        plt.suptitle("Homoscedasticity Test - Residuals vs Predicted Values", fontsize=18)
        plt.tight_layout()
        plt.show()
```

Homoscedasticity Test - Residuals vs Predicted Values



OBSERVATION

For Train data

- Pattern in Residuals:
 - The residuals is randomly scattered around the horizontal axis (y = 0), indicating that the model captures the underlying data pattern well.
 - There is a clear pattern (such as a curve), it may indicate non-linearity.
- Homoscedasticity:

■ The residuals gather in certain areas, it indicates there may be heteroscedasticity, meaning the variance of the errors is not constant, which may affect the model's predictions.

For Test data

- Pattern in Residuals:
 - The residuals is randomly scattered around the horizontal axis (y = 0), indicating that the model captures the underlying data pattern well.
 - There is a clear pattern (such as a curve), it may indicate non-linearity.
- Homoscedasticity:
 - The spread of residuals remains consistent across all levels of predicted values indicating homoscedasticity, meaning the variance of the errors constant.

```
In [ ]: # Import the Library
        from statsmodels.stats.diagnostic import het_goldfeldquandt
        # Perform Goldfeld-Quandt test for train data
        gq_test = het_goldfeldquandt(y_train, X_train)
        f_statistic = gq_test[0]
        p_value = gq_test[1]
        # Print the results
        print("Goldfeld-Quandt test F-statistic:", f_statistic)
        print("Goldfeld-Quandt test p-value:", p_value)
        # Interpretation
        if p_value > 0.05:
            print("There is no strong evidence of heteroscedasticity. Homoscedasticity is validated.")
        else:
            print("There is strong evidence of heteroscedasticity.")
       Goldfeld-Quandt test F-statistic: 1.0029318358365393
       Goldfeld-Quandt test p-value: 0.4918984144278115
       There is no strong evidence of heteroscedasticity. Homoscedasticity is validated.
```

```
In []: # Import the library
from statsmodels.stats.diagnostic import het_goldfeldquandt

# Perform Goldfeld-Quandt test for test data
gq_test = het_goldfeldquandt(y_test, X_test)
f_statistic = gq_test[0]
p_value = gq_test[1]

# Print the results
print("Goldfeld-Quandt test F-statistic:", f_statistic)
print("Goldfeld-Quandt test p-value:", p_value)

# Interpretation
if p_value > 0.05:
    print("There is no strong evidence of heteroscedasticity. Homoscedasticity is validated.")
else:
    print("There is strong evidence of heteroscedasticity.")
```

Goldfeld-Quandt test F-statistic: 1.2278089004946002 Goldfeld-Quandt test p-value: 0.2519350659398718 There is no strong evidence of heteroscedasticity. Homoscedasticity is validated.

OBSERVATION

• Since the p-value is greater than 0.05, we conclude that there is no strong evidence of heteroscedasticity in the dataset. This implies that the assumption of constant variance of residuals (homoscedasticity) holds for our linear regression model.

9.3.5) Normality of residuals

Normality of residuals refers to the assumption that the residuals (or errors) in a statistical model are normally distributed. Residuals are the differences between the observed values and the predicted values from the model.


The assumption of normality is important in many statistical analyses because it allows for the application of certain statistical tests and the validity of confidence intervals and hypothesis tests. When residuals are normally distributed, it implies that the errors are random, unbiased, and have consistent variability.


To check for the normality of residuals, you can follow these steps:

- **Residual Histogram:** Create a histogram of the residuals and visually inspect whether the shape of the histogram resembles a bell-shaped curve. If the majority of the residuals are clustered around the mean with a symmetric distribution, it suggests normality.
- **Q-Q Plot** (Quantile-Quantile Plot): This plot compares the quantiles of the residuals against the quantiles of a theoretical normal distribution. If the points in the Q-Q plot are reasonably close to the diagonal line, it indicates that the residuals are normally distributed. Deviations from the line may suggest departures from normality.

• Shapiro-Wilk Test: This is a statistical test that checks the null hypothesis that the residuals are normally distributed. The Shapiro-Wilk test calculates a test statistic and provides a p-value. If the p-value is greater than the chosen significance level (e.g., 0.05), it suggests that the residuals follow a normal distribution. However, this test may not be reliable for large sample sizes. Anderson-Darling or Jarque_Bera can also be done as data size increases

```
In [ ]: # Calculate residuals
        train_residuals = y_train - y_pred_train
        test_residuals = y_test - y_pred_test
        # Plot histogram and Q-Q plot for train residuals
        plt.figure(figsize=(15, 5))
        plt.subplot(1, 2, 1)
        sns.histplot(train_residuals, kde=True, color='green')
        plt.title('Histogram of Train Residuals')
        plt.subplot(1, 2, 2)
        stats.probplot(train_residuals, dist="norm", plot=plt)
        plt.title('Q-Q Plot of Train Residuals')
        plt.suptitle('Normality of Residuals - Train Data', fontsize=18)
        plt.show()
        # Plot histogram and Q-Q plot for test residuals
        plt.figure(figsize=(15, 5))
        plt.subplot(1, 2, 1)
        sns.histplot(test_residuals, kde=True, color='green')
        plt.title('Histogram of Test Residuals')
        plt.subplot(1, 2, 2)
        stats.probplot(test_residuals, dist="norm", plot=plt)
        plt.title('Q-Q Plot of Test Residuals')
        plt.suptitle('Normality of Residuals - Test Data', fontsize=18)
        plt.show()
```





```
In [ ]: # Shapiro-Wilk Test for normality
import scipy.stats as stats
```

```
# HO: Data is Gaussian
        # Ha: Data is not Gaussian
        shapiro_train = stats.shapiro(train_residuals)
        shapiro_test = stats.shapiro(test_residuals)
        print(f'Shapiro-Wilk Test for Train Residuals: Statistic={shapiro_train.statistic}, p-value={shapiro_train.pvalue}')
        if shapiro_train.pvalue > 0.05:
            print("Train Data is Gaussian (fail to reject H0)")
        else:
            print("Train Data is not Gaussian (reject H0)")
        print()
        print(f'Shapiro-Wilk Test for Test Residuals: Statistic={shapiro_test.statistic}, p-value={shapiro_test.pvalue}')
        if shapiro_test.pvalue > 0.05:
            print("Test Data is Gaussian (fail to reject H0)")
            print("Test Data is not Gaussian (reject H0)")
       Shapiro-Wilk Test for Train Residuals: Statistic=0.9291010499000549, p-value=7.73526370994454e-13
       Train Data is not Gaussian (reject H0)
       Shapiro-Wilk Test for Test Residuals: Statistic=0.9178698658943176, p-value=1.0869382094824687e-05
       Test Data is not Gaussian (reject H0)
In [ ]: # Anderson-Darling Test for normality
        from statsmodels.stats.diagnostic import normal_ad
        # HO: Data is Gaussian
        # Ha: Data is not Gaussian
        anderson train = normal ad(train residuals)
        anderson_test = normal_ad(test_residuals)
        print(f'Anderson-Darling Test for Train Residuals: Statistic={anderson_train[0]}, p-value={anderson_train[1]}')
        if anderson_train[1] > 0.05:
            print("Train Data is Gaussian (fail to reject H0)")
        else:
            print("Train Data is not Gaussian (reject H0)")
        print()
        print(f'Anderson-Darling Test for Test Residuals: Statistic={anderson_test[0]}, p-value={anderson_test[1]}')
        if anderson_test[1] > 0.05:
            print("Test Data is Gaussian (fail to reject H0)")
        else:
            print("Test Data is not Gaussian (reject H0)")
       Anderson-Darling Test for Train Residuals: Statistic=7.357345909096807, p-value=5.303849926227439e-18
       Train Data is not Gaussian (reject H0)
```

Anderson-Darling Test for Test Residuals: Statistic=2.1093362831425537, p-value=2.128546023452642e-05 Test Data is not Gaussian (reject H0)

OBSERVATION

train_scores = []

• The Shapiro-Wilk test and Anderson-Darling Test indicates that the residuals from the training and testing data do not follow a Gaussian distribution.

9.4) Polynomial Regression 💹

9.4.1) Find the best degree

```
In []: # Function for Adj. R2 Score

def adj_r(r_sq,X,Y):
    adj_r1 = (1 - ((1-r_sq)*(len(Y)-1)) / (len(Y)-X.shape[1]-1))
    return adj_r1

def r2_score(y,y_):
    num = np.sum((y-y_)**2)
    denom = np.sum((y- y.mean())**2)
    score = (1- num/denom)
    return score

In []: # Creating a pipeline
    from sklearn.pipeline import make_pipeline
    from sklearn.preprocessing import PolynomialFeatures
    from sklearn.metrics import mean_squared_error
    degrees = 4
```


```
test_scores = []
train_loss = []
test_loss = []
for degree in range(1, degrees):
    # Putting the classes like PolynomialFeatures(), StandardScaler(), LinearRegression() into a pipeline
    polyreg_scaled = make_pipeline(PolynomialFeatures(degree), StandardScaler(), LinearRegression())
    polyreg_scaled.fit(X_train, y_train)
    # Calculate R2 Score for train and test data
    train_score = polyreg_scaled.score(X_train, y_train) # R2 TRAIN
    test_score = polyreg_scaled.score(X_test, y_test) # R2 TEST
    # Calculate Adj. R2 Score for train and test data
    train_scores.append(adj_r(train_score,X_train,y_train))
    test_scores.append(adj_r(test_score,X_test,y_test))
    # Calculate the y_pred for train and test data
    output1 = polyreg_scaled.predict(X_train)
    output2 = polyreg_scaled.predict(X_test)
    # Calculate the MSE for train and test data
    train_loss.append(mean_squared_error(y_train,output1)) # MSE train
    test_loss.append(mean_squared_error(y_test,output2)) # MSE test
```

```
In []: # Plote
fig, axes = plt.subplots(1, 2, figsize=(12, 3))

axes[0].plot(list(range(1, degrees)), train_scores, label="train")
axes[0].plot(list(range(1, degrees)), test_scores, label="test")
axes[0].set_xlabel("degree")
axes[0].set_ylabel("Adj.R-score")

axes[1].plot(list(range(1, degrees)), train_loss, label="train")
axes[1].plot(list(range(1, degrees)), test_loss, label="test")
axes[1].legend(loc='upper left')
axes[1].set_xlabel("degree")
axes[1].set_ylabel("MSE")

plt.show()
```



- As we go in higher Degree, the model test performance drop significantly Which clearly indicates Overfitting
- The Test score is maximum at degree 2, There for the **best polynomial degree is 2**

```
In [ ]: rate_list = list(range(1, degrees))
In [ ]: # Best degree
    index = np.argmax(test_scores)
    best_degree = rate_list[index]
    best_degree
Out[ ]: 2
```

OBSERVATION

• The best degree is 2, We are going to use this degree for further evaluation.

9.4.2) Polynomial Regression implementation

```
In []: # Polynomial Regression

# Transform the features into polynomial features
from sklearn.preprocessing import PolynomialFeatures
degree = 2
```

```
poly = PolynomialFeatures(degree=degree)
        X_train_poly = poly.fit_transform(X_train)
        X_test_poly = poly.transform(X_test)
In [ ]: # Shape of X_train and X_test
        X_train.shape, X_test.shape
Out[]: ((400, 7), (100, 7))
In [ ]: # Polynomial features been created
        X_train_poly.shape, X_test_poly.shape
Out[]: ((400, 36), (100, 36))
In [ ]: # Standardize the polynomial features
        from sklearn.preprocessing import StandardScaler
        scaler = StandardScaler()
        X_train_poly_scaled = scaler.fit_transform(X_train_poly)
        X_test_poly_scaled = scaler.transform(X_test_poly)
In [ ]: # Import the required library
        from sklearn.linear_model import LinearRegression
        # Initialize the Linear Regression model
        Polynomial_Reg_model = LinearRegression()
        # Train the model
        Polynomial_Reg_model.fit(X_train_poly_scaled, y_train)
        ▼ LinearRegression
Out[ ]:
        LinearRegression()
In [ ]: # Predicting values for the test data
        y_pred_train_poly = Polynomial_Reg_model.predict(X_train_poly_scaled)
        y_pred_test_poly = Polynomial_Reg_model.predict(X_test_poly_scaled)
In [ ]: # Performance of Polynomial Regression
        print("Performance of Polynomial Regression")
        print("-"*36)
        print("Performance of Train data")
        print("-"*26)
        evaluate_model_performance(y_train, y_pred_train_poly, Polynomial_Reg_model)
        print("Performance of Test data")
        print("-"*26)
        evaluate_model_performance(y_test, y_pred_test_poly, Polynomial_Reg_model)
       Performance of Polynomial Regression
       Performance of Train data
       MSE: 0.0032
       MAE: 0.04
       RMSE: 0.0569
       R-squared: 0.84
       Adjusted R-squared: 0.82
       Performance of Test data
       MSE: 0.0035
       MAE: 0.0406
       RMSE: 0.0596
       R-squared: 0.83
       Adjusted R-squared: 0.73
```

- Polynomial Regression provides a slight improvement over Linear Regression in terms of training data performance with an R-squared value of 0.84.
- However, there is a drop in the Adjusted R-squared value for the test data (0.73), indicating potential overfitting. The model fits the training data well but doesn't generalize as well to unseen data.

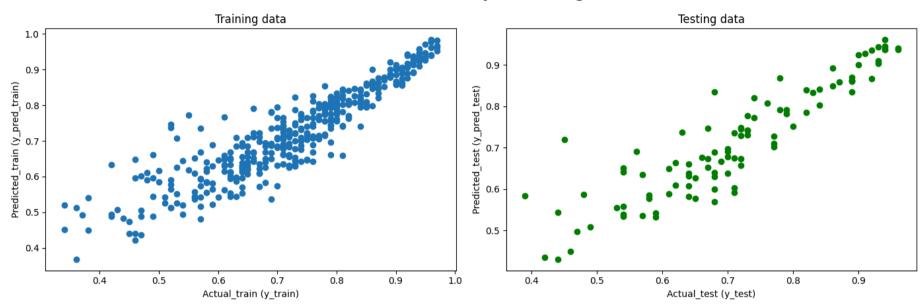
```
In []: # Actual vs Predicted
plt.figure(figsize=(14, 5))

# Actual vs Predicted Plot
plt.subplot(1, 2, 1)
plt.scatter(y_train, y_pred_train_poly)
plt.xlabel('Actual_train (y_train)')
plt.ylabel('Predicted_train (y_pred_train)')
plt.title('Training data')
```

```
# Actual vs Predicted Plot
plt.subplot(1, 2, 2)
plt.scatter(y_test, y_pred_test_poly, color="green")
plt.xlabel('Actual_test (y_test)')
plt.ylabel('Predicted_test (y_pred_test)')
plt.title('Testing data')

plt.suptitle("Actual vs Predicted - Polynomial Regression", fontsize = 18)
plt.tight_layout()
plt.show()
```

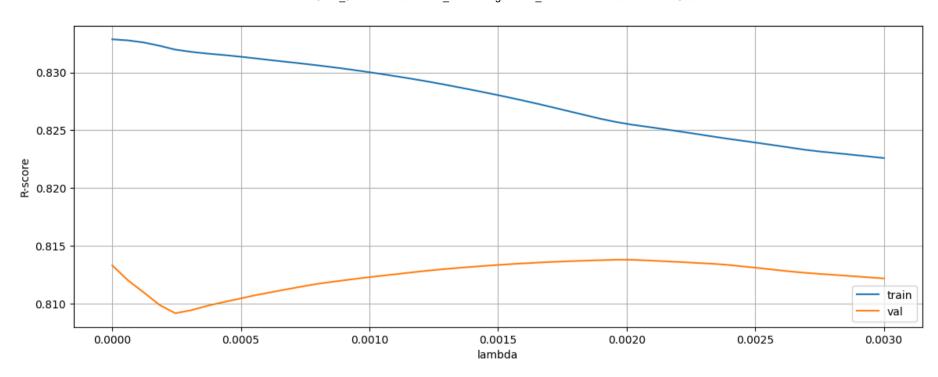
Actual vs Predicted - Polynomial Regression



9.5) Regularization

9.5.1) Laso regression model

```
In [ ]: # Hyperparameter Tuning: find the best regularization strength
        from sklearn.linear_model import Lasso, Ridge
        # To find best Lambda
        degree = 2 # is best
        train_scores = []
        test_scores = []
        rate_list = np.linspace(0,0.003,50)
        for rate in rate_list:
          # Creating pipeline()
          polyreg_scaled = make_pipeline(PolynomialFeatures(2), StandardScaler(), Lasso(alpha=rate))
          polyreg_scaled.fit(X_train, y_train)
          # Calculate R2 Score for train and test data
          train_score = polyreg_scaled.score(X_train, y_train)
          test_score = polyreg_scaled.score(X_test, y_test)
          # Calculate Adj. R2 Score for train and test data
          train_scores.append(adj_r(train_score,X_train,y_train))
          test_scores.append(adj_r(test_score,X_test,y_test))
        # Plote
        plt.figure(figsize=(14, 5))
        plt.plot(rate_list, train_scores, label="train")
        plt.plot(rate_list, test_scores, label="val")
        plt.legend(loc='lower right')
        plt.xlabel("lambda")
        plt.ylabel("R-score")
        plt.grid()
        plt.show()
```



```
In [ ]: # Best Lambda (or) alpha
index = np.argmax(test_scores)
best_lambda = rate_list[index]
best_lambda
```

```
Out[]: 0.0019591836734693877
```

```
In [ ]: # Final Lasso model
        # degree 2 and Lambda :0.0019591836734693877
        final_lasso_model_pipe = make_pipeline(PolynomialFeatures(2), StandardScaler(), Lasso(alpha=best_lambda))
        final_lasso_model_pipe.fit(X_train, y_train)
        # Predicting values for the train and test data
        y_pred_train_lasso = final_lasso_model_pipe.predict(X_train)
        y_pred_test_lasso = final_lasso_model_pipe.predict(X_test)
        # Performance of Lasso Regression
        print("Performance of Lasso Regression")
        print("-"*36)
        # Metrix for train and test data
        print("Performance of Train data")
        print("-"*26)
        evaluate_model_performance(y_train, y_pred_train_lasso, final_lasso_model_pipe)
        print()
        print("Performance of Test data")
        print("-"*26)
        evaluate_model_performance(y_test, y_pred_test_lasso, final_lasso_model_pipe)
```

- Lasso Regression performs similarly to Linear Regression with a slight improvement in the training data R-squared value (0.83).
- The model maintains a consistent Adjusted R-squared value of 0.81 on the test data, suggesting a good balance between model complexity and generalization.

```
In []: # Actual vs Predicted
plt.figure(figsize=(14, 5))

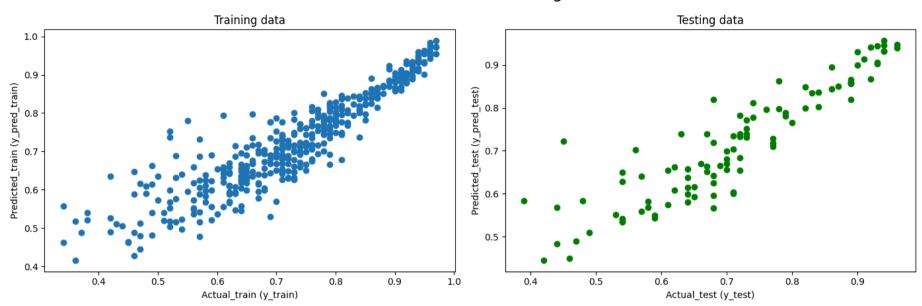
# Actual vs Predicted Plot
plt.subplot(1, 2, 1)
plt.scatter(y_train, y_pred_train_lasso)
plt.xlabel('Actual_train (y_train)')
plt.ylabel('Predicted_train (y_pred_train)')
plt.title('Training data')

# Actual vs Predicted Plot
plt.subplot(1, 2, 2)
plt.scatter(y_test, y_pred_test_lasso, color="green")
```

```
plt.xlabel('Actual_test (y_test)')
plt.ylabel('Predicted_test (y_pred_test)')
plt.title('Testing data')

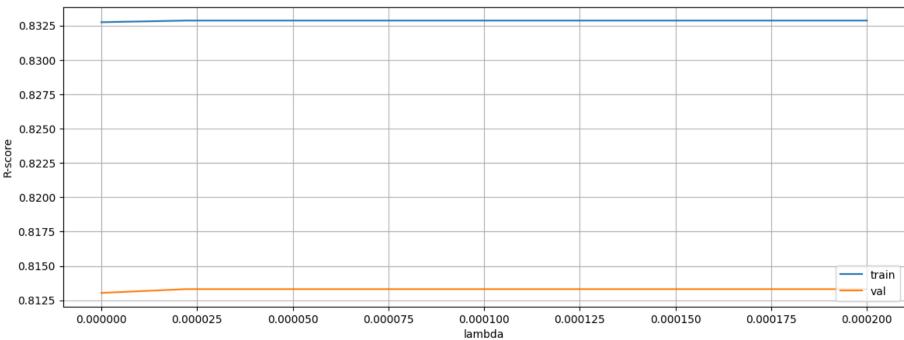
plt.suptitle("Actual vs Predicted - Lasso Regression", fontsize = 18)
plt.tight_layout()
plt.show()
```

Actual vs Predicted - Lasso Regression



9.5.2) Ridge regression model

```
In [ ]: # Hyperparameter Tuning: find the best regularization strength
        from sklearn.linear_model import Lasso, Ridge
        # To find best Lambda
        degree = 2 # is best
        train_scores = []
        test_scores = []
        rate_list = np.linspace(0,0.0002,10)
        for rate in rate_list:
          # Creating pipeline()
          polyreg_scaled = make_pipeline(PolynomialFeatures(2), StandardScaler(), Ridge(alpha=rate))
          polyreg_scaled.fit(X_train, y_train)
          # Calculate R2 Score for train and test data
          train_score = polyreg_scaled.score(X_train, y_train)
          test_score = polyreg_scaled.score(X_test, y_test)
          # Calculate Adj. R2 Score for train and test data
          train_scores.append(adj_r(train_score,X_train,y_train))
          test_scores.append(adj_r(test_score,X_test,y_test))
        # Plote
        plt.figure(figsize=(14, 5))
        plt.plot(rate_list, train_scores, label="train")
        plt.plot(rate_list, test_scores, label="val")
        plt.legend(loc='lower right')
        plt.xlabel("lambda")
        plt.ylabel("R-score")
        plt.grid()
        plt.show()
```

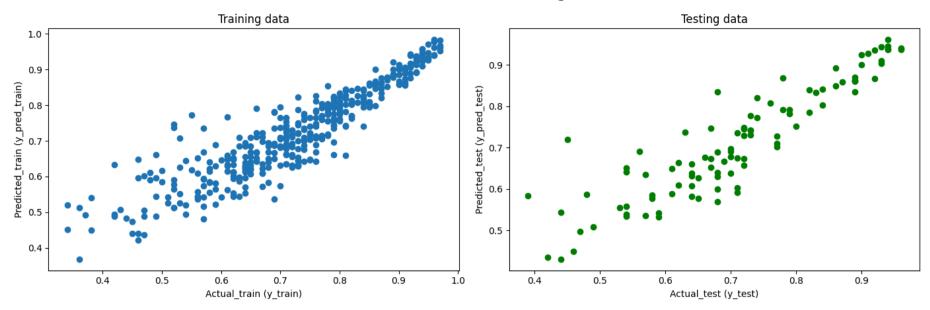


```
In [ ]: # Best Lambda (or) alpha
        index = np.argmax(test_scores)
        best_lambda_Ridge = rate_list[index]
        best_lambda_Ridge
Out[]: 2.2222222222223e-05
In [ ]: # Final Lasso model
        # degree: 2 and Lambda :2.2222222222222e-05
        final_ridge_model_pipe = make_pipeline(PolynomialFeatures(2), StandardScaler(), Ridge(alpha=best_lambda_Ridge))
        final_ridge_model_pipe.fit(X_train, y_train)
        # Predicting values for the train and test data
        y_pred_train_ridge = final_ridge_model_pipe.predict(X_train)
        y_pred_test_ridge = final_ridge_model_pipe.predict(X_test)
        # Performance of Ridge Regression
        print("Performance of Ridge Regression")
        print("-"*36)
        # Metrix for train and test data
        print("Performance of Train data")
        print("-"*26)
        evaluate_model_performance(y_train, y_pred_train_ridge, final_ridge_model_pipe)
        print()
        print("Performance of Test data")
        print("-"*26)
        evaluate_model_performance(y_test, y_pred_test_ridge, final_ridge_model_pipe)
       Performance of Ridge Regression
       Performance of Train data
       MSE: 0.0032
       MAE: 0.04
       RMSE: 0.0569
       R-squared: 0.84
       Adjusted R-squared: 0.83
       Performance of Test data
       MSE: 0.0035
       MAE: 0.0406
       RMSE: 0.0596
       R-squared: 0.83
       Adjusted R-squared: 0.81
```

- Ridge Regression shows the best performance in terms of training R-squared (0.84) and maintains a good performance on test data with an Adjusted R-squared of 0.81.
- The model demonstrates good generalization with low MSE, MAE, and RMSE values.


```
In [ ]: # Actual vs Predicted
        plt.figure(figsize=(14, 5))
        # Actual vs Predicted Plot
        plt.subplot(1, 2, 1)
        plt.scatter(y_train, y_pred_train_ridge)
        plt.xlabel('Actual_train (y_train)')
        plt.ylabel('Predicted_train (y_pred_train)')
        plt.title('Training data')
        # Actual vs Predicted Plot
        plt.subplot(1, 2, 2)
        plt.scatter(y_test, y_pred_test_ridge, color="green")
        plt.xlabel('Actual_test (y_test)')
        plt.ylabel('Predicted_test (y_pred_test)')
        plt.title('Testing data')
        plt.suptitle("Actual vs Predicted - Lasso Regression", fontsize = 18)
        plt.tight_layout()
        plt.show()
```

Actual vs Predicted - Lasso Regression



9.5.3) Elastic Net regression model

```
In [ ]: # Hyperparameter Tuning: find the best regularization strength
        from sklearn.linear_model import ElasticNet
        # To find best Lambda
        degree = 2 # is best
        train_scores = []
        test_scores = []
        rate_list = np.linspace(0,0.006,10)
        for rate in rate_list:
          # Creating pipeline()
          polyreg_scaled = make_pipeline(PolynomialFeatures(2), StandardScaler(), ElasticNet(alpha=rate))
          polyreg_scaled.fit(X_train, y_train)
          # Calculate R2 Score for train and test data
          train_score = polyreg_scaled.score(X_train, y_train)
          test_score = polyreg_scaled.score(X_test, y_test)
          # Calculate Adj. R2 Score for train and test data
          train_scores.append(adj_r(train_score,X_train,y_train))
          test_scores.append(adj_r(test_score,X_test,y_test))
        # Plote
        plt.figure(figsize=(14, 5))
        plt.plot(rate_list, train_scores, label="train")
        plt.plot(rate_list, test_scores, label="val")
        plt.legend(loc='lower right')
        plt.xlabel("lambda")
        plt.ylabel("R-score")
        plt.grid()
        plt.show()
```



```
In []: # Best Lambda (or) alpha
index = np.argmax(test_scores)
best_lambda_ElasticNet = rate_list[index]
best_lambda_ElasticNet
```

Out[]: 0.004

```
In [ ]: # Final ElasticNet model
        # degree: 2 and Lambda: 0.004
        final_ElasticNet_model_pipe = make_pipeline(PolynomialFeatures(2), StandardScaler(), Ridge(alpha=best_lambda_ElasticNet))
        final_ElasticNet_model_pipe.fit(X_train, y_train)
        # Predicting values for the train and test data
        y_pred_train_ElasticNet = final_ElasticNet_model_pipe.predict(X_train)
        y_pred_test_ElasticNet = final_ElasticNet_model_pipe.predict(X_test)
        # Performance of ElasticNet Regression
        print("Performance of ElasticNet Regression")
        print("-"*36)
        # Metrix for train and test data
        print("Performance of Train data")
        print("-"*26)
        evaluate_model_performance(y_train, y_pred_train_ElasticNet, final_ElasticNet_model_pipe)
        print()
        print("Performance of Test data")
        print("-"*26)
        evaluate_model_performance(y_test, y_pred_test_ElasticNet, final_ElasticNet_model_pipe)
       Performance of ElasticNet Regression
```

OBSERVATION

- ElasticNet Regression also performs well with the highest training R-squared (0.84) and a consistent Adjusted R-squared (0.81) on the test data.
- The performance metrics are similar to Ridge Regression, indicating a strong balance between bias and variance.

9.6) Model Conclusion @

Ridge Regression and ElasticNet Regression are recommended for their strong performance and balance between model complexity and generalization. Both models have demonstrated consistent metrics across training and test data, indicating their robustness for predicting the chance of admission.

10) Inference 🐇

• University Preference:

■ The majority of applicants (32%) prefer universities with a rating of 3, followed by ratings of 2 (25%) and 4 (21%). This indicates a significant preference for mid-tier universities among applicants.

• Research Experience:

• 56% of the applicants have research experience, suggesting that having research experience is a common trait among those applying for graduate programs.

• Skewed Admission Chances:

The chances of admission are negatively skewed, meaning more applicants have higher probabilities of being admitted, indicating overall strong applicant profiles.

• Key Factors for Admission:

 CGPA has the highest correlation with the chance of admission, followed by GRE and TOEFL scores. This highlights the importance of academic performance in the admission process.

• University Ratings and Admission:

 Higher university ratings are associated with higher chances of admission, emphasizing the impact of the university's prestige on admission decisions.

• Impact of Research:

Applicants with research experience have a higher chance of admission, especially noticeable in universities with the highest rating

• Performance of Models:

Polynomial, Lasso, Ridge, and ElasticNet regressions all show similar performance, slightly better than simple linear regression, indicating that more complex models can capture the nuances of the data better.

• Model Metrics:

All models have high R-squared values (around 0.82 to 0.84) for train data, indicating good fit, and slightly lower for test data, showing consistency but room for improvement in generalization.

• Importance of CGPA:

Among all features, CGPA carries the highest weight, followed by GRE and TOEFL scores, underscoring the critical role of academic excellence.

• Model Comparison:

• While Polynomial and Ridge regressions provide the best performance, the differences across models are marginal. This suggests that multiple approaches can be viable for predicting admission chances.

11) Recommendations 🤝

• Focus on Academic Excellence:

• Encourage students to maintain a high CGPA, as it significantly impacts their admission chances.

• **GRE and TOEFL Preparation:**

Offer resources and support for GRE and TOEFL preparation since these scores are crucial for admission.

• Promote Research Opportunities:

■ Facilitate more research opportunities for students, as research experience notably increases admission probabilities.

• Guidance on University Selection:

• Provide tailored advice to students on selecting universities based on their profiles, focusing on those with mid to high ratings.

• Application Strategy:

Develop strategies for students to apply to a balanced mix of universities (ratings 2 to 4) to maximize their admission chances.

• Highlighting Research Impact:

■ Emphasize the benefits of research experience in marketing materials and counseling sessions.

• Improving Admission Predictors:

Continuously refine and validate admission prediction models using updated data to enhance accuracy.

• Resource Allocation:

Allocate more resources to support areas identified as critical, such as GRE/TOEFL prep and research opportunities.

Ву

SANKARAN MEENAKSHI SUNDARAM

DSML July 2023